

```
In [1]: import torch
import torch.nn.functional
from torch import nn

from labml import experiment, tracker
from labml.configs import option
from labml_helpers.train_valid import BatchIndex
from labml_nn.distillation.large import LargeModel
from labml_nn.distillation.small import SmallModel
from labml_nn.experiments.cifar10 import CIFAR10Configs
```

```
In [2]: class Configs(CIFAR10Configs):
    model: SmallModel
    large: LargeModel
    kl_div_loss = nn.KLDivLoss(log_target=True)
    loss_func = nn.CrossEntropyLoss()
    temperature: float = 5.
    soft_targets_weight: float = 100.
    label_loss_weight: float = 0.5
```

```

In [3]: def step(self, batch: any, batch_idx: BatchIndex):
        self.model.train(self.mode.is_train)
        self.large.eval()
        data, target = batch[0].to(self.device), batch[1].to(self.device)
        if self.mode.is_train:
            tracker.add_global_step(len(data))
        with torch.no_grad():
            large_logits = self.large(data)
            output = self.model(data)
            soft_targets = nn.functional.log_softmax(large_logits / self.temperature, dim=-1)
            soft_prob = nn.functional.log_softmax(output / self.temperature, dim=-1)
            soft_targets_loss = self.kl_div_loss(soft_prob, soft_targets)
            label_loss = self.loss_func(output, target)
            loss = self.soft_targets_weight * soft_targets_loss + self.label_loss_weight * label_loss
            tracker.add({"loss.kl_div.": soft_targets_loss,
                        "loss.nll": label_loss,
                        "loss.": loss})
        self.accuracy(output, target)
        self.accuracy.track()
        if self.mode.is_train:
            loss.backward()
            self.optimizer.step()
            if batch_idx.is_last:
                tracker.add('model', self.model)
            self.optimizer.zero_grad()
        tracker.save()

```

```

In [4]: @option(Configs.large)
        def _large_model(c: Configs):
            return LargeModel().to(c.device)
        @option(Configs.model)
        def _small_student_model(c: Configs):
            return SmallModel().to(c.device)

```

```
In [5]: def get_saved_model(run_uuid: str, checkpoint: int):
        from labml_nn.distillation.large import Configs as LargeConfigs
        experiment.evaluate()
        conf = LargeConfigs()
        experiment.configs(conf, experiment.load_configs(run_uuid))
        experiment.add_pytorch_models({'model': conf.model})
        experiment.load(run_uuid, checkpoint)
        experiment.start()
        return conf.model
```

```
In [6]: def main(run_uuid: str, checkpoint: int):
        large_model = get_saved_model(run_uuid, checkpoint)
        experiment.create(name='distillation', comment='cifar10')
        conf = Configs()
        conf.large = large_model
        experiment.configs(conf, {
            'optimizer.optimizer': 'Adam',
            'optimizer.learning_rate': 2.5e-4,
            'model': '_small_student_model',
        })
        experiment.add_pytorch_models({'model': conf.model})
        experiment.load(None, None)
        with experiment.start():
            conf.run()
```

```
In [7]: main('d46cd53edaec11eb93c38d6538aee7d6', 1_000_000)
```

LABML WARNING

.labml.yaml config file could not be found. Looking in path: /Users/jianyiyang/Desktop/work/deepLearning/summerClass/final/distillation example

LABML WARNING

Not a valid git repository: /Users/jianyiyang/Desktop/work/deepLearning/summerClass/final/distillation example

LABML ERROR

Couldn't find a previous run to load configurations: d46cd53edaec11eb93c38d6538aee7d6

Prepare model...

Prepare device.device...

Prepare device.device_info...[DONE] 2.69ms

Prepare device.device...[DONE] 6.46ms

Prepare model...[DONE] 110.67ms

Couldn't find a previous run

Notebook Experiment: 83b07f462d4311ee9310acde48001122

[clean]: ""

LABML WARNING

Not a valid git repository: /Users/jianyiyang/Desktop/work/deepLearning/summerClass/final/distillation example

Prepare model...[DONE] 10.28ms

No labml server url specified. Please start a labml server and specify the URL. Docs: <https://github.com>

om/labmlai/labml/tree/master/app

distillation: 83f598e22d4311ee9310acde48001122

cifar10

[clean]: ""

~/labml/configs.yaml does not exist. Creating /Users/jianyiyang/.labml/configs.yaml

Initialize...

Prepare mode...[DONE] 3.25ms

Initialize...[DONE] 77.61ms

Prepare validator...

Prepare valid_loader...

Prepare valid_dataset...

Prepare dataset_transforms...[DONE] 3.33ms

Prepare valid_dataset...[DONE] 63,081.73ms

Prepare valid_loader...[DONE] 63,177.20ms

Prepare validator...[DONE] 63,277.03ms

Prepare trainer...

Prepare train_loader...

Prepare train_dataset...[DONE] 725.73ms

Prepare train_loader...[DONE] 819.43ms

Prepare trainer...[DONE] 846.29ms

Prepare training_loop...

Prepare loop_count...[DONE] 46.19ms

Prepare training_loop...[DONE] 229.86ms

0: Train: ... 0ms 0ms 0:00m/ 0:00m

Prepare optimizer...[DONE] 3.94ms

Prepare optimizer.optimizer...

Prepare optimizer.weight_decay_obj...[DONE] 4.14ms

Prepare optimizer.optimizer...[DONE] 10.60ms

50,000: Train: 100% 245,426ms Valid: 100% 11,634ms loss.train: 1.051353 accuracy.train: 0.578660
loss.valid: 0.91503 accuracy.valid: 0.608100 321,307ms 0:05m/ 0:48m

100,000: Train: 100% 233,280ms Valid: 100% 11,383ms loss.train: 0.858824 accuracy.train: 0.730740
loss.valid: 0.738977 accuracy.valid: 0.725500 284,795ms 0:09m/ 0:37m

150,000: Train: 100% 241,689ms Valid: 100% 11,771ms loss.train: 0.730364 accuracy.train: 0.7914
20 loss.valid: 0.676379 accuracy.valid: 0.765300 275,545ms 0:14m/ 0:31m

200,000: Train: 100% 244,112ms Valid: 100% 11,960ms loss.train: 0.626952 accuracy.train: 0.8314
00 loss.valid: 0.656888 accuracy.valid: 0.774500 279,666ms 0:18m/ 0:27m


```
In [8]: print(LargeModel())
```



```

LargeModel(
  (layers): Sequential(
    (0): Dropout(p=0.1, inplace=False)
    (1): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (2): BatchNorm()
    (3): ReLU(inplace=True)
    (4): Dropout(p=0.1, inplace=False)
    (5): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (6): BatchNorm()
    (7): ReLU(inplace=True)
    (8): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (9): Dropout(p=0.1, inplace=False)
    (10): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): BatchNorm()
    (12): ReLU(inplace=True)
    (13): Dropout(p=0.1, inplace=False)
    (14): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (15): BatchNorm()
    (16): ReLU(inplace=True)
    (17): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (18): Dropout(p=0.1, inplace=False)
    (19): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (20): BatchNorm()
    (21): ReLU(inplace=True)
    (22): Dropout(p=0.1, inplace=False)
    (23): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (24): BatchNorm()
    (25): ReLU(inplace=True)
    (26): Dropout(p=0.1, inplace=False)
    (27): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (28): BatchNorm()
    (29): ReLU(inplace=True)
    (30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (31): Dropout(p=0.1, inplace=False)
    (32): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (33): BatchNorm()
    (34): ReLU(inplace=True)
    (35): Dropout(p=0.1, inplace=False)
    (36): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (37): BatchNorm()
    (38): ReLU(inplace=True)
    (39): Dropout(p=0.1, inplace=False)
    (40): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  )
)

```

```
(41): BatchNorm()  
(42): ReLU(inplace=True)  
(43): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
(44): Dropout(p=0.1, inplace=False)  
(45): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
(46): BatchNorm()  
(47): ReLU(inplace=True)  
(48): Dropout(p=0.1, inplace=False)  
(49): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
(50): BatchNorm()  
(51): ReLU(inplace=True)  
(52): Dropout(p=0.1, inplace=False)  
(53): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
(54): BatchNorm()  
(55): ReLU(inplace=True)  
(56): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
)  
(fc): Linear(in_features=512, out_features=10, bias=True)  
)
```

In []: